

AAanalysis

v1.1 Cheat Sheet

Sequence → physicochemical scales → interpretable features → explainable ML → biological mechanism

[pip install aaanalysis](#) · [aaanalysis.readthedocs.io](#)

AAanalysis is a Python framework for interpretable, sequence-based protein prediction. It turns sequences into physicochemical features (CPP), trains explainable models, and traces every prediction back to a residue × property × group comparison — robust for small datasets. **v1.1** extends the core feature engine beyond physicochemical scales to PLM embeddings and protein structure.

Install · Import · Load

core + [pro]

```
# Python >= 3.11
pip install aaanalysis          # core
pip install 'aaanalysis[pro]'  # SHAP, FIMO, Bio

import numpy as np
import matplotlib.pyplot as plt
import aaanalysis as aa
df_seq = aa.load_dataset(name='DOM_GSEC') # γ-secretase
labels = df_seq['label'].to_list()
df_scales = aa.load_scales()
```

The Golden Workflow

canonical pipeline

- LOAD** load_dataset · load_scales → df_seq · df_scales
- PARTS** SequenceFeature.get_df_parts → df_parts
- FEATURES** Part × Split × Scale · CPP.run → df_feat
- MODEL** TreeModel.fit · dPULearn.fit → feat_importance · labels_
- EXPLAIN** CPPPlot.feature_map · ShapModel → figure · feat_impact

Prediction Task Levels

task → setup

Residue AA_*
unit: sliding window (aa_window_size)
ref: non-site windows / shuffled background

positional

Domain DOM_*
unit: part-set jmd_n · tmd · jmd_c (from tmd_start / tmd_stop)
ref: labelled A vs B groups

both

Protein SEQ_*
unit: whole chain (composition)
ref: labelled groups / composition-matched background

compositional

Output types: classification · regression · ranking · explanation

Groups: 1: positives 0: negatives 0: reliable negatives 2: unlabeled
Label 0 is a curated **negative** in labelled data, or a dPULearn-inferred **reliable negative** drawn from the unlabeled (2) pool.

Sequence Anatomy

the TMD model

TMD Target Middle Domain — the central segment of interest (e.g. transmembrane stretch, binding region).

JMD Juxta Middle Domain — the flanks adjoining the TMD (jmd_n on the N-side, jmd_c on the C-side).

JMD-N TMD JMD-C
0 tmd_start tmd_stop len(seq)
JMD widths set globally: aa.options['jmd_n_len'] · ['jmd_c_len'].

CPP Feature Concept

Part × Split × Scale

PART where on the sequence

tmd · jmd_n · jmd_c · tmd_jmd · jmd_n_tmd_n · tmd_c_jmd_c

SPLIT how to read the part

Segment — contiguous · Pattern — sparse pairs · PeriodicPattern — i, i+3/4

SCALE which physicochemical property

AAontology (~600 scales) · hydrophobicity · charge · helix propensity

TMD A I I G L M V G G V V I
Segment(1,4) ■ ■ ■ ■ ■ ■ ■ ■ ■ ■
Pattern(N,1,4,8) ■ ■ ■ ■ ■ ■ ■ ■ ■ ■
PeriodicPattern ■ ■ ■ ■ ■ ■ ■ ■ ■ ■
Simplified from Breimann25a (Suppl. Fig. 1C ↗)

TMD × Segment × hydrophobicity → membrane insertion

JMD × Pattern × net charge → electrostatic recognition

TMD × PeriodicPattern × helix → α-helical interface

CPP Strategies

via split_kws

Compositional ≈ sequence/protein-level

one whole-part average (composition-like, position-agnostic)

```
split_kws = sf.get_split_kws(
    split_types="Segment",
    n_split_max=1)
cpp = aa.CPP(df_parts=df_parts, split_kws=split_kws)
```

Positional ≈ residue-/region-level

sub-segments and/or patterns resolved to positions

```
split_kws = sf.get_split_kws(
    split_types=["Segment", "Pattern", "PeriodicPattern"],
    n_split_max=5,
    steps_pattern=[3, 4],
    steps_periodicpattern=[3, 4])
cpp = aa.CPP(df_parts=df_parts, split_kws=split_kws)
```

Domain level uses both. → CPP strategies: see the CPP tutorial (docs).

Which Module Should I Use?

intent → module

Explore sequence patterns / composition

AAlogo

Sample reference windows (if negatives are missing)

AAWindowSampler

Reduce redundant amino acid scales

AAclust

Discover discriminative physicochemical features

CPP

Train with positives + unlabeled data

dPULearn

Train an interpretable classifier

TreeModel

Explain a prediction (per feature / sample)

ShapModel [pro]

Data & Preparation

load · clean

Load benchmark sequences

load_dataset(name) → df_seq

Load AAontology scales

load_scales() → df_scales

Load precomputed features

load_features(name) → df_feat

Read / write FASTA

read_fasta(file) → df_seq

Cluster redundant homologs

filter_seq(df_seq) → df_clust [pro]

Sequence Analysis

logos · motifs

Position-specific logo

AAlogo().get_df_logo(df_parts) → df_logo

Sample reference windows

AAWindowSampler().sample_*(df_seq)

Pairwise sequence similarity

comp_seq_sim(df_seq) [pro]

Scan motifs (FIMO / MEME)

scan_motif(df_seq, pwm) → df_hits [pro]

Feature Engineering

parts · CPP · scales

SequenceFeature → sf

sf = aa.SequenceFeature()

· split sequence into parts

sf.get_df_parts(df_seq) → df_parts

· assemble feature matrix X

sf.feature_matrix(df_feat, df_parts) → X

Discover discriminative features

CPP(df_parts).run(labels) → df_feat *

Sweep CPP configs (grid)

CPPGrid().run(...) · .eval() → ranked configs

Simplify → interpretable scales

CPP.simplify(df_feat, labels) → df_feat

Reduce redundant scales

AAclust().fit(X) [Wrapper]

Drop correlated features

NumericalFeature().filter_correlation(X)

Feature Preprocessing

one-hot · PLM · structure · PTM

Encode sequences (one-hot / int)

SequencePreprocessor().encode_*(seqs) → X

PLM embeddings v1.1

EmbeddingPreprocessor().encode(...) → dict_num

Structure / DSSP / PAE v1.1

StructurePreprocessor().encode_dssp(...) → dict_num [pro]

PTM / site annotations v1.1

AnnotationPreprocessor().encode(...) → dict_num [pro]

Combine sources v1.1

combine_dict_nums([...]) → dict_num

Numerical CPP v1.1

CPP(df_parts).run_num(dict_num_parts, labels) → df_feat

Modeling & Explainability

Train with positives + unlabeled data

dPULearn().fit(X, labels) [Wrapper]

Train + RFE + MC importance

TreeModel().fit(X, labels) [Wrapper]

Per-feature / sample SHAP impact

ShapModel().fit(X, labels) [pro]

Metrics & Plotting

utilities

Adjusted AUC (class imbalance)

comp_auc_adjusted(X, labels)

BIC score · KL divergence

comp_bic_score(X, labels) · comp_kld

Per-protein / detection (v1.1)

comp_per_protein_ap · comp_detection_metrics

Plot style, fonts & standalone legend

plot_settings(font_scale) · plot_legend(ax)

Protein Design (to be extended)

In-silico point mutations v1.1

AAmut · AAMutPlot


Sequence-design libraries v1.1

SeqMut · SeqMutPlot

AAlogo — see the data

dataset at a glance

```
import numpy as np, matplotlib.pyplot as plt, aaanalysis as aa
df_seq = aa.load_dataset(name='DOM_GSEC') # γ-secretase
labels = list(df_seq['label']); df_scales = aa.load_scales()
sf = aa.SequenceFeature()
df_parts = sf.get_df_parts(df_seq=df_seq,
    list_parts=['tmd', 'jmd_n', 'jmd_c'])
aa.plot_settings(font_scale=0.7)
# aal_kws builds df_logo + bits bar for you
aa.AAlogoPlot().single_logo(
    aal_kws=dict(df_parts=df_parts, labels=labels,
        label_test=1, tmd_len=20),
    name_data='Test set: substrates')
plt.tight_layout(); plt.show()
```



AAlogoPlot.single_logo · per-position enrichment

CPP — feature

top feature · REF vs TEST

```
# default parts + a redundancy-reduced set of 100 scales
df_parts = sf.get_df_parts(df_seq=df_seq)
df_scales = aa.AAclust().select_scales(
    df_scales=df_scales, n_clusters=100)
cpp = aa.CPP(df_parts=df_parts, df_scales=df_scales)
df_feat = cpp.run(labels=labels, n_filter=100)
X = sf.feature_matrix(df_feat['feature'], df_parts)
tm = aa.TreeModel(); tm.fit(X, labels=labels)
df_feat = tm.add_feat_importance(df_feat=df_feat, sort=True)
cpp_plot = aa.CPPPlot(); aa.plot_settings()
# distribution of the top feature (feat_rank=1 of the sorted
df_feat)
cpp_plot.feature(feature=df_feat, feat_rank=1, df_seq=df_seq,
    labels=labels, name_test='substrates', name_ref='non-subs.')
plt.tight_layout(); plt.show()
```


▼ figure below · left

CPP — ranking


top features · effect + importance

```
# same df_feat — rank the top discriminative features
aa.plot_settings(font_scale=0.6)
cpp_plot.ranking(df_feat=df_feat, n_top=15, rank=True,
    name_test='substrates', name_ref='non-subs.')
plt.tight_layout(); plt.show()
```

▼ figure below · right



CPPPlot.feature · top feature, REF vs TEST




CPPPlot.ranking · top-15 features


CPP — feature map

group level · full vs simplified

```
# global Part × Split × Scale map — all AAontology scales
cpp_plot = aa.CPPPlot(); aa.plot_settings(font_scale=0.65)
cpp_plot.feature_map(df_feat=df_feat)
# CPP.simplify → fewer, interpretable correlated scales
df_feat = cpp.simplify(df_feat=df_feat, labels=labels)
cpp_plot.feature_map(df_feat=df_feat)
plt.tight_layout(); plt.show()
```



CPPPlot.feature_map · all scales




CPPPlot.feature_map · simplified


ShapModel — explain a prediction

sample level · adva

```
# advanced: per-sample explanation (fuzzy labeling demo)
# fuzzy labeling: APP's label is its soft prediction score (0.6,
not 1)
i = list(df_seq['entry']).index('P05067') # APP
y = [float(v) for v in labels]; y[i] = 0.6
sm = aa.ShapModel(); sm.fit(X, labels=y, fuzzy_labeling=True)
df_feat = sm.add_feat_impact(df_feat=df_feat,
    sample_positions=i, names='APP')
args_seq = {k + '_seq': v for k, v in sf.get_df_parts(
    df_seq=df_seq).loc['P05067'].to_dict().items()}
ka = dict(col_imp='feat_impact APP', shap_plot=True, **args_seq)
cpp_plot.profile(df_feat=df_feat, **ka)
cpp_plot.feature_map(df_feat=df_feat, name_test='APP', **ka)
plt.tight_layout(); plt.show()
```



CPPPlot.profile · SHAP



CPPPlot.feature_map · SHAP

AAclust — clusters

scale reduction · Wrapper

```
aac = aa.AAclust()
aac.select_scales(df_scales, n_clusters=10)
aac.medoid_names_ # 10 reduced scales (labels_ also set)

aac_plot = aa.AAclustPlot()
aac_plot.centers(np.array(df_scales).T, labels=aac.labels_)
plt.tight_layout(); plt.show()
```

▼ figure below · left


dPULearn — PCA

reliable negatives · Wrapper


```
# DOM_GSEC ships 1/0 — encode as PU labels: 1 = positive, 2 =
unlabeled
labels_pu = [1 if x == 1 else 2 for x in labels]
n_pos = sum(np.array(labels_pu) == 1)
dpul = aa.dPULearn()
dpul.fit(X=X, labels=labels_pu, n_unl_to_neg=n_pos // 2)
df_pu = dpul.df_pu_ # 1 pos · 0 rel-neg · 2 unl
```

```
dpul_plot = aa.dPULearnPlot()
dpul_plot.pca(df_pu=df_pu, labels=dpul.labels_)
plt.tight_layout(); plt.show()
```

▼ figure below · right



AAclustPlot.centers · cluster scale profiles



dPULearnPlot.pca · reliable negatives

AAWindowSampler

build reference windows

```
# Reference windows around sites when you lack negatives:
aaws = aa.AAWindowSampler()
# SAME proteins · window 9 (odd) → PTM / single-residue site
df_same = aaws.sample_same_protein(df_seq, n=100, window_size=9)
# DIFFERENT proteins · window 10 (even) → cleavage bond
df_diff = aaws.sample_different_protein(df_seq, n=100,
    window_size=10)
# SYNTHETIC — AA-frequency priors (null background)
df_syn = aaws.sample_synthetic(df_seq, n=100,
    generator='global_freq')
```

What are you predicting?

per residue / site → AA_* · odd/even window · parts = window
per domain / region → DOM_* · TMD model · parts = jmd_n-tmd:jmd_c
whole protein → SEQ_* · composition · whole chain

What labels do you have?

labeled 0 / 1 → CPP → ML model
positives + unlabeled (1 / 2) → CPP → dPULearn → ML model
no negatives at all → AAWindowSampler → CPP → ML model

What is your learning task?

classify → TreeModel
regress → get_labels_quantile / tiered → TreeModel
multi-class → get_labels_ovr / ovo
cluster · PCA → AAclust · dPULearn

Which explainability do you need?

global (group level) → CPP → TreeModel → CPPPlot
local (per protein) → CPP → ShapModel → CPPPlot

- Gotchas
- things that bite
- Labels: **1/0** = supervised (pos/neg). **dPULearn needs 1/2** (pos/unlabeled) and outputs **0 = reliable-negative**.
 - load_dataset(name, n=N)** returns **2N** rows (N per class) — count classes via df_seq['label'].
 - Compositional vs positional is not a flag — it **emerges from split_kws**.
 - Reproducibility: layered seeds — seed= ► random_state= ► options['random_state'] ► default.
 - DOM_*** parts need tmd_start/tmd_stop in df_seq; **[pro]** features need pip install 'aaanalysis[pro]'.
 - TMD = Target** Middle Domain — generalized from **Transmembrane** domain (Breimann25a) to any central segment; JMD-N / JMD-C are its flanks.

Glossary

Feature (CPP) — (Part × Split × Scale) — the atomic, residue-grounded, interpretable unit of CPP.

Part — Named segment used as feature input: tmd, jmd_n, jmd_c, tmd_jmd, jmd_n_tmd_n, tmd_c_jmd_c.

Split — How a scale is read across a part: Segment (contiguous), Pattern (sparse), PeriodicPattern (i, i+3/4).

Scale — AA → ℝ mapping. AAontology ships ~600 curated scales in two-level categories.

AAontology — Two-level scale taxonomy; CPP uses its categories to organize and rank features.

Design Principles

- Explicit over implicit — DataFrames everywhere
- Wrappers (.fit / .predict / .eval) set trailing *_attributes after fit
- Biological interpretability is first-class
- Small-data robust and reproducible (layered seeds)

Data Objects

shapes & columns

df_seq	entry · sequence · label · tmd_start · tmd_stop
df_parts	one column per part: tmd · jmd_n · jmd_c · ...
df_feat	feature · abs_auc · mean_dif · p_val · positions · scale · category
X	feature matrix (samples × features) from sf.feature_matrix
dict_num	{entry: ndarray (L×D)} — numerical per-residue values (v1.1)

Class · abbr ↔ Plot Class

mirrored API

CLASS	ABBR	PLOT CLASS	KIND
SequencePreprocessor	sp	—	
EmbeddingPreprocessor	ep	—	
StructurePreprocessor	pro stp	—	
AnnotationPreprocessor	pro ap	—	
AAlogo	aal	AAlogoPlot	
AAWindowSampler	aaws	—	
SequenceFeature	sf	—	
CPP	cpp	CPPPlot	
AAclust	aac	AAclustPlot	Wrapper
NumericalFeature	nf	—	
dPULearn	dpu_l	dPULearnPlot	Wrapper
TreeModel	tm	—	Wrapper
ShapModel	pro sm	—	Wrapper
AAMut	aamut	AAMutPlot	
SeqMut	seqmut	SeqMutPlot	

CPP — Comparative Physicochemical Profiling — discovers ranked Part × Split × Scale features.

Test vs reference group — The A-vs-B contrast CPP profiles: a feature's mean_dif is test — reference (name_test / name_ref in CPPPlot).

Compositional vs positional — How split_kws resolves locality: a whole-part average (compositional) vs sub-region/position-resolved (positional).

Numerical CPP (pseudo-scale) v1.1 — CPP generalizes from AA→scale lookup to any per-residue tensor — PLM · structure · PTM — each a pseudo-scale via CPP.run_num.

Feature importance vs impact — Two explainability axes: importance = unsigned, group-level (TreeModel); impact = signed, per-sample (ShapModel, shap_plot).

aa.options

system-level settings

```
aa.options['random_state'] = 42
aa.options['verbose'] = True
aa.options['n_jobs'] = -1 # all cores (None = auto)
aa.options['allow_multiprocessing'] = True

# TMD model — JMD flank widths
aa.options['jmd_n_len'] = 10
aa.options['jmd_c_len'] = 10

# plot labels & system-level scales
aa.options['name_tmd'] = 'P5-P5' # e.g. cleavage site
aa.options['df_scales'] = my_scales
```

How to Cite

if you use AAanalysis

AAclust [Breimann24a]
Breimann & Frishman (2024a), AAclust: k-optimized clustering for selecting redundancy-reduced sets of amino acid scales
Bioinformatics Advances ↗

AAontology [Breimann24b]
Breimann et al. (2024b), AAontology: An ontology of amino acid scales for interpretable machine learning
Journal of Molecular Biology ↗

CPP & dPULearn [Breimann25a]
Breimann & Kamp et al. (2025), Charting γ-secretase substrates by explainable AI
Nature Communications ↗

canonical vocabulary · CONTEXT.md

Reducing features — Four distinct ops: redundancy reduction (AAclust scales) · feature pruning · selection (RFE) · simplification (CPP.simplify → interpretable scales).

PU labels — dPULearn input: 1 = positive, 2 = unlabeled. Output: 1 / 0 (reliable-negative) / 2.

Wrapper class — sklearn-style class — .fit / .predict / .eval, sets trailing *_attributes after fit.

Plot class — *Plot mirror of an analytical class — same arguments, visualization only.